

Abhishek K

9760832448 · abhishekrestha06@gmail.com · github.com/mky62 · abhishekk-phi.vercel.app

SUMMARY

Full-stack developer specializing in AI-integrated products — RAG pipelines, LLM streaming, vector search, and agentic workflows. Ships end-to-end: schema, APIs, React UI, and cloud deployment. Looking for a high-velocity team building at the intersection of AI and developer tooling.

PROJECTS

devnarrate · *Next.js* · *Node.js* · *PostgreSQL* · *Prisma* · *Redis* · *Pinecone* · *OpenRouter* · *Tiptap*

github.com/mky62/devnarrate

- Built an AI-assisted developer publishing platform — GitHub OAuth, public profiles, repository showcase, technical blogging — shipped solo from schema to production.
- Designed a RAG pipeline: chunked and embedded repository code via OpenRouter, stored vectors in Pinecone, retrieved relevant context, and streamed AI-generated technical drafts to the editor — cut content creation time by ~60%.
- Redis caching on GitHub API reduced hot-path calls by ~70%; PostgreSQL/Prisma schema covers users, sessions, posts, likes, and contribution analytics.

cortex · *Turborepo* · *Next.js* · *Convex* · *Vapi* · *Gemini* · *OpenRouter* · *Clerk*

github.com/mky62/cortex

- Architected a multi-tenant AI customer support SaaS in a Turborepo monorepo — operator dashboard and embeddable widget sharing a shadcn/ui component library, with Clerk org-scoped auth.
- Built a RAG knowledge-base pipeline: ingested PDFs/images via Convex storage, extracted text with Gemini multimodal API, generated OpenRouter embeddings in org-isolated namespaces — improved AI answer relevance by ~45%.
- Integrated Vapi voice support end-to-end: AWS Secrets Manager credential storage, assistant discovery, and widget-side WebRTC calling with real-time Convex agent threads and human handoff routing.

reviewbot · *Next.js* · *FastAPI* · *OpenAI GPT-4o* · *GitHub API* · *Docker* · *Redis*

github.com/mky62/reviewbot

- AI GitHub PR reviewer: triggers on pull requests, retrieves diffs via GitHub REST API, and posts structured review comments using GPT-4o with custom prompts enforcing team style guides.
- Multi-pass pipeline — security scan (OWASP heuristics), logic correctness, and style — reduced review turnaround from hours to under 2 minutes.
- Containerized with Docker, deployed on Railway with GitHub App webhooks; Redis dedup prevents double-reviews on force-pushes.

voicequery · *Next.js* · *Whisper API* · *LangChain* · *PostgreSQL* · *pgvector* · *Vercel AI SDK* github.com/mky62/voicequery

- Voice-to-SQL natural language interface: Whisper transcribes speech, LangChain constructs schema-aware prompts, app returns live DB results with no SQL knowledge required.
- Semantic query caching with pgvector: embedded past queries and served near-duplicate results from cache — cut LLM API calls by ~55% on repeated question patterns.
- Streamed responses token-by-token via Vercel AI SDK, achieving sub-second time-to-first-token over tabular data.

SKILLS

AI / ML: RAG pipelines, LLM integration (OpenAI GPT-4o, Gemini, OpenRouter), LangChain, vector search (Pinecone, pgvector), Whisper, Hugging Face, prompt engineering, agentic workflows, AI streaming (Vercel AI SDK)

Frontend: React, Next.js, React Native, TailwindCSS, shadcn/ui

Backend: Node.js, FastAPI, REST APIs, WebSockets, Convex, Vapi

Databases: PostgreSQL, MongoDB, Firebase, Prisma, Redis

DevOps: Docker, GitHub Actions (CI/CD), Vercel, Railway, AWS, Turborepo

Languages: TypeScript, JavaScript, Python, Java

EDUCATION
